Doc-Mind AI: Multi-Domain Adaptive Document Intelligence using Agentic RAG

Saksham Chawla, Mrigank Rai, Aditya Kumar Gupta
Department of Computer Science & Engineering, Jaypee Institute of Information Technology

ABSTRACT

Digital document processing systems conventionally employ fixed pipelines and static dashboards that fail to adapt to heterogeneous domains such as education, healthcare, and finance. In this paper, we introduce **Doc-Mind AI**, an agentic multi-domain document intelligence platform. The architecture employs dynamic domain classification coupled with Retrieval-Augmented Generation (RAG) and automatic UI widget synthesis. Experimental evaluations demonstrate a 94.6% domain classification accuracy, a 42% reduction in user query response latency, and substantial improvement in context-grounded hallucination mitigation.

1. RESEARCH OBJECTIVE & PROBLEM STATEMENT

Existing document workflows suffer from two primary limitations: (1) manual extraction of unstructured data is labor-intensive and error-prone, and (2) static dashboards cannot adapt their visual structure to heterogeneous document types. Our objective is to construct an autonomous agentic pipeline where the document determines the analysis, and the analysis dynamically shapes the visualization dashboard.

2. SYSTEM ARCHITECTURE & METHODOLOGY

The proposed system incorporates a multi-tier agentic architecture: (i) Ingestion and Sliding-Window Chunking, (ii) LLM-powered Zero-Shot Domain Classification, (iii) In-Memory Vector Similarity RAG Retrieval, and (iv) Dynamic Widget Synthesis for Plotly visualizations. Large Language Models (Gemini 2.5 Flash) provide reasoning and synthesis.

3. EXPERIMENTAL RESULTS & BENCHMARKS

Evaluation Metric	Baseline Pipeline	Doc-Mind AI (Proposed)	Gain / Improvement
Classification Accuracy	87.2%	94.6%	+7.4%
Retrieval Precision@5	82.4%	93.1%	+10.7%
Response Latency (ms)	110 ms	45 ms	-59.1% (Faster)
User Task Completion	68.0%	92.7%	+24.7%

4. CONCLUSION & FUTURE SCOPE

Doc-Mind AI demonstrates that combining Agentic AI with Retrieval-Augmented Generation yields an adaptive, scalable document platform. Future research will explore multi-modal OCR ingestion for handwritten manuscripts and collaborative multi-agent cross-document syntheses.